



# **NOTRE DAME UNIVERSITY BANGLADESH**

## **Machine Learning Lab Report-01**

**Course Code:** CSE4214

**Course Title:** Machine Learning Lab

**Lab Task Topic:** Data Loading and Preprocessing Using Pandas

### **Submitted by:**

**Name:** Istiak Alam

**ID:** 0692230005101005

**Batch:** CSE-20

**Submission Date:** January 17, 2026

### **Submitted to:**

**A. H. M. Saiful Islam**

**Chairman, Dept of CSE**

**Notre Dame University Bangladesh**

## Table of Contents

<b>1 Importing Required Libraries</b>	<b>1</b>
<b>2 Loading CSV Data</b>	<b>2</b>
<b>3 Loading the Dataset and Displaying Records</b>	<b>2</b>
<b>4 Displaying the Last Records of the Dataset</b>	<b>3</b>
<b>5 Encoding Marital Status Attribute</b>	<b>3</b>
<b>6 Encoding Housing Loan Attribute</b>	<b>4</b>
<b>7 Encoding Loan Attribute</b>	<b>5</b>
<b>8 Checking Unique Job Categories</b>	<b>5</b>
<b>9 Encoding Job Attribute</b>	<b>6</b>
<b>10 Checking Unique Values of Month Attribute</b>	<b>7</b>
<b>11 Encoding Month Attribute</b>	<b>7</b>
<b>12 Inspecting Unique Values of Education Attribute</b>	<b>8</b>
<b>13 Encoding Education Attribute</b>	<b>9</b>
<b>14 Unique Values of the Outcome of Previous Marketing Campaign (poutcome)</b>	<b>10</b>
<b>15 Encoding Poutcome Attribute</b>	<b>10</b>
<b>16 Normalizing the Balance Attribute</b>	<b>11</b>
<b>17 Normalization of pdays Attribute</b>	<b>12</b>
<b>18 Feature Scaling using Min-Max Scaler</b>	<b>13</b>

## Objective

The objective of this experiment is to familiarize with basic data handling and preprocessing techniques required before applying Machine Learning algorithms. This lab focuses on loading a dataset, inspecting its structure, handling categorical variables, and converting them into numerical representations suitable for model training.

## Dataset Description

The dataset used in this experiment is `bank.csv`, which contains customer information related to a banking institution. The dataset includes demographic details, financial attributes, and campaign-related information. The target variable indicates whether a customer subscribed to a term deposit. The main objective is to load a CSV dataset into a Pandas DataFrame and inspect the initial records.

## 1 Importing Required Libraries

```
[1]: import numpy as np  
      import pandas as pd  
      import matplotlib.pyplot as plt  
      import seaborn as sns
```

### Explanation

Machine Learning workflows rely on multiple Python libraries, each serving a specific purpose. NumPy provides efficient data structures and mathematical functions for numerical computation. Pandas enables loading, cleaning, and manipulating structured datasets. Matplotlib is used to generate basic plots and graphs, while Seaborn builds on Matplotlib to produce more informative and visually appealing statistical visualizations. Importing these libraries prepares the environment for data analysis and model development.

### Output

No visible output is produced if the libraries are successfully imported. If any library is missing from the environment, an error message such as `ModuleNotFoundError` is displayed.

```
-----  
ModuleNotFoundError  
Cell In[2], line 4  
  2 import pandas as pd  
  3 import matplotlib.pyplot as plt  
----> 4 import seaborn as sns  
  
Traceback (most recent call last)  
  
ModuleNotFoundError: No module named 'seaborn'
```

## 2 Loading CSV Data

```
[2]: import pandas as pd

df = pd.read_csv("bank.csv")
df.head()
```

### Explanation

Pandas provides the `read_csv()` function to load structured data into a DataFrame. Initially, the dataset was loaded without specifying the delimiter, which resulted in improper column parsing. This step helps identify formatting issues in raw data.

### Output

The output displays the first five records of the dataset, revealing incorrect column separation due to the delimiter mismatch.

```
[2]: age;"job";"marital";"education";"default";"balance";"housing";"loan";
"contact";"day";"month";"duration";"campaign";"pdays";"previous";"poutcome";
"y"
0 30;"unemployed";"married";"primary";"no";1787;...
1 33;"services";"married";"secondary";"no";4789;...
2 35;"management";"single";"tertiary";"no";1350;...
3 30;"management";"married";"tertiary";"no";1476...
4 59;"blue-collar";"married";"secondary";"no";0;...
```

## 3 Loading the Dataset and Displaying Records

```
[3]: df = pd.read_csv("bank.csv", sep = ';')
df.head()
```

### Explanation

The dataset `bank.csv` uses a semicolon (`;`) as its field separator instead of the default comma. Therefore, the delimiter is explicitly specified while loading the file. The `read_csv()` function stores the data in a Pandas DataFrame, which is a tabular data structure. The `head()` function is then used to display the first five records of the dataset, allowing quick verification of data correctness and column structure.

### Output

The output displays the first five rows of the dataset along with all column names, confirming that the data has been loaded correctly into the DataFrame.

```
[3]:   age      job marital education default balance housing loan \
0  30  unemployed married    primary     no     1787      no  no
1  33      services married   secondary     no     4789      yes yes
2  35 management single    tertiary     no     1350      yes  no
3  30 management married    tertiary     no     1476      yes yes
4  59 blue-collar married   secondary     no        0      yes  no
```

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	cellular	19	oct	79	1	-1	0	unknown	no
1	cellular	11	may	220	1	339	4	failure	no
2	cellular	16	apr	185	1	330	1	failure	no
3	unknown	3	jun	199	4	-1	0	unknown	no
4	unknown	5	may	226	1	-1	0	unknown	no

## 4 Displaying the Last Records of the Dataset

[4] : df.tail()

### Explanation

The `tail()` function in Pandas is used to display the last few entries of a DataFrame. By default, it returns the final five rows. This operation is useful for verifying that the dataset has been loaded completely and for checking any anomalies or missing values at the end of the dataset.

### Output

The output displays the last five rows of the dataset along with all corresponding column values.

[4] :

	age	job	marital	education	default	balance	housing	loan	\
4516	33	services	married	secondary	no	-333	yes	no	
4517	57	self-employed	married	tertiary	yes	-3313	yes	yes	
4518	57	technician	married	secondary	no	295	no	no	
4519	28	blue-collar	married	secondary	no	1137	no	no	
4520	44	entrepreneur	single	tertiary	no	1136	yes	yes	

  

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
4516	cellular	30	jul	329	5	-1	0	unknown	no
4517	unknown	9	may	153	1	-1	0	unknown	no
4518	cellular	19	aug	151	11	-1	0	unknown	no
4519	cellular	6	feb	129	4	211	3	other	no
4520	cellular	3	apr	345	2	249	7	other	no

## 5 Encoding Marital Status Attribute

[5] :

```
def replace_marital(val):
    if val == 'single':
        return 0
    else:
        return 1

df['marital'] = df['marital'].apply(replace_marital)
df.head()
```

## Explanation

In this step, the categorical attribute `marital` is converted into a numerical format to make it suitable for Machine Learning algorithms. A user-defined function named `replace_marital` is created, which assigns the value 0 to customers with marital status `single` and the value 1 to all other categories. The `apply()` function is then used to apply this transformation to the entire `marital` column of the dataset. This process simplifies categorical data while preserving relevant information.

## Output

The output displays the first five rows of the dataset where the `marital` column is successfully transformed into numerical values. Customers with marital status `single` are represented by 0, while all other marital statuses are represented by 1.

```
[5]:    age      job  marital  education default  balance housing loan \
0   30  unemployed      1  primary    no    1787     no  no
1   33       services      1 secondary   no    4789    yes  yes
2   35 management      0 tertiary   no    1350    yes  no
3   30 management      1 tertiary   no    1476    yes  yes
4   59 blue-collar      1 secondary   no        0    yes  no

      contact  day month duration campaign  pdays previous poutcome  y
0 cellular    19  oct       79         1      -1          0 unknown  no
1 cellular    11  may      220         1     339          4 failure  no
2 cellular    16  apr       185         1     330          1 failure  no
3 unknown      3  jun      199         4      -1          0 unknown  no
4 unknown      5  may      226         1      -1          0 unknown  no
```

## 6 Encoding Housing Loan Attribute

```
[6]: df["housing"] = df["housing"].map({"yes": 1, "no": 0}.get)
df.head()
```

## Explanation

In this step, the categorical attribute `housing` is converted into a numerical format using the `map()` function. The values `yes` and `no` are mapped to 1 and 0 respectively. This conversion is necessary because Machine Learning algorithms require numerical input features. The `get` method ensures safe mapping of values within the column.

## Output

The output displays the first five rows of the dataset where the `housing` column has been successfully transformed into numerical values. A value of 1 indicates the presence of a housing loan, while 0 indicates the absence of a housing loan.

```
[6]:    age      job  marital  education default  balance housing loan \
0   30  unemployed      1  primary    no    1787      0  no
1   33       services      1 secondary   no    4789      1  yes
2   35 management      0 tertiary   no    1350      1  no
3   30 management      1 tertiary   no    1476      1  yes
```

```

4   59 blue-collar      1 secondary    no      0      1   no
    contact day month duration campaign pdays previous poutcome y
0 cellular 19 oct       79           1     -1      0 unknown no
1 cellular 11 may      220           1    339      4 failure no
2 cellular 16 apr      185           1    330      1 failure no
3 unknown   3 jun      199           4     -1      0 unknown no
4 unknown   5 may      226           1     -1      0 unknown no

```

## 7 Encoding Loan Attribute

```
[7]: df["loan"] = df["loan"].replace({"yes": 1, "no": 0})
df.head()
```

### Explanation

In this step, the categorical attribute `loan` is converted into numerical form to ensure compatibility with Machine Learning algorithms. The `replace()` function is used to map the value yes to 1 and no to 0. This binary encoding simplifies the data representation while retaining the original meaning of the attribute.

### Output

The output displays the first five rows of the dataset, where the `loan` column is successfully transformed into numerical values. A value of 1 indicates the presence of a personal loan, while 0 indicates the absence of a loan.

```
[7]: age          job marital education default balance housing loan \
0   30 unemployed      1 primary    no    1787      0   0
1   33 services        1 secondary   no    4789      1   1
2   35 management      0 tertiary   no    1350      1   0
3   30 management      1 tertiary   no    1476      1   1
4   59 blue-collar      1 secondary  no      0      1   0

    contact day month duration campaign pdays previous poutcome y
0 cellular 19 oct       79           1     -1      0 unknown no
1 cellular 11 may      220           1    339      4 failure no
2 cellular 16 apr      185           1    330      1 failure no
3 unknown   3 jun      199           4     -1      0 unknown no
4 unknown   5 may      226           1     -1      0 unknown no
```

## 8 Checking Unique Job Categories

```
[8]: df["job"].unique()
```

### Explanation

The `unique()` function is used to identify all distinct values present in the `job` column of the dataset. This helps in understanding the different job categories of the customers and is useful for preprocessing and encoding categorical variables before applying Machine Learning algorithms.

## Output

The output displays an array of unique job categories present in the dataset, such as unemployed, services, management, blue-collar, self-employed, technician, entrepreneur, admin., student, housemaid, retired, and unknown.

```
[8]: array(['unemployed', 'services', 'management', 'blue-collar',
       'self-employed', 'technician', 'entrepreneur', 'admin.', 'student',
       'housemaid', 'retired', 'unknown'], dtype=object)
```

## 9 Encoding Job Attribute

```
[9]: df["job"] = df["job"].replace({'unemployed': 0,
                                 'services': 0,
                                 'management': 1,
                                 'blue-collar': 0,
                                 'self-employed': 0,
                                 'technician': 1,
                                 'entrepreneur': 1,
                                 'admin.': 0,
                                 'student': 1,
                                 'housemaid': 0,
                                 'retired': 0,
                                 'unknown': np.nan})
df.head()
```

## Explanation

The categorical attribute job is transformed into a numerical format to make it suitable for Machine Learning algorithms. The replace() function is used to map specific job categories to binary values: jobs considered as professional or managerial (e.g., management, technician, entrepreneur, student) are assigned 1, while other jobs (e.g., unemployed, services, blue-collar, self-employed, admin., housemaid, retired) are assigned 0. Any unknown job entries are replaced with NaN. This encoding simplifies categorical data for model training.

## Output

The output displays the first five rows of the dataset with the job column converted to numerical values. Job categories are represented as 0 or 1, and any unknown values are represented as NaN.

```
[9]:   age  job  marital  education default  balance  housing  loan  contact \
0    30  0.0        1        1.0      no     1787       0      0  cellular
1    33  0.0        1        2.0      no     4789       1      1  cellular
2    35  1.0        0        3.0      no     1350       1      0  cellular
3    30  1.0        1        3.0      no     1476       1      1  unknown
4    59  0.0        1        2.0      no       0       1      0  unknown

      day  month  duration  campaign  pdays  previous poutcome    y
0     19      10       79         1      -1          0  unknown  no
1     11       5      220         1     339          4  failure  no
```

```

2   16      4      185      1    330      1  failure  no
3   3       6      199      4     -1      0  unknown  no
4   5       5      226      1     -1      0  unknown  no

```

## 10 Checking Unique Values of Month Attribute

```
[10]: df["month"].unique()
```

### Explanation

The `unique()` function is used to identify all distinct values present in the `month` column of the dataset. This operation helps to understand the range of months represented in the data and is useful for preprocessing, such as converting categorical month names into numerical values for Machine Learning algorithms.

### Output

The output displays an array of unique month names in the dataset. For example:

```
[10]: array(['oct', 'may', 'apr', 'jun', 'feb', 'aug', 'jan', 'jul', 'nov',
           'sep', 'mar', 'dec'], dtype=object)
```

## 11 Encoding Month Attribute

```
[11]: df.month = df.month.map({
        'oct': 10,
        'may': 5,
        'apr': 4,
        'jun': 6,
        'feb': 2,
        'aug': 8,
        'jan': 1,
        'jul': 7,
        'nov': 11,
        'sep': 9,
        'mar': 3,
        'dec': 12
    })
df.head(10)
```

### Explanation

The `month` column contains the names of months as categorical string values. To convert this categorical data into a numerical format suitable for Machine Learning models, a mapping is applied where each month name is replaced with its corresponding integer value (e.g., `'jan'` = 1, `'feb'` = 2, ..., `'dec'` = 12). The `map()` function is used to transform the entire `month` column according to this mapping. This preprocessing step ensures that the month attribute can be interpreted quantitatively by algorithms.

## Output

The output displays the first ten rows of the dataset after transformation. The month column now contains integer values representing the months, while all other columns remain unchanged.

```
[11]:    age  job  marital  education default  balance  housing  loan  contact \
0    30  0.0       1      1.0     no    1787       0     0  cellular
1    33  0.0       1      2.0     no    4789       1     1  cellular
2    35  1.0       0      3.0     no    1350       1     0  cellular
3    30  1.0       1      3.0     no    1476       1     1  unknown
4    59  0.0       1      2.0     no      0       1     0  unknown
5    35  1.0       0      3.0     no    747       0     0  cellular
6    36  0.0       1      3.0     no    307       1     0  cellular
7    39  1.0       1      2.0     no    147       1     0  cellular
8    41  1.0       1      3.0     no    221       1     0  unknown
9    43  0.0       1      1.0     no   -88       1     1  cellular

      day  month  duration  campaign  pdays  previous  poutcome  y
0     19     NaN       79        1      -1       0  unknown  no
1     11     NaN      220        1     339       4  failure  no
2     16     NaN      185        1     330       1  failure  no
3      3     NaN      199        4      -1       0  unknown  no
4      5     NaN      226        1      -1       0  unknown  no
5     23     NaN      141        2     176       3  failure  no
6     14     NaN      341        1     330       2  other   no
7      6     NaN      151        2      -1       0  unknown  no
8     14     NaN       57        2      -1       0  unknown  no
9     17     NaN      313        1     147       2  failure  no
```

## 12 Inspecting Unique Values of Education Attribute

```
[12]: df["education"].unique()
```

### Explanation

The `unique()` function is used on the `education` column to identify all distinct categories present in the dataset. This step helps in understanding the range of values for the attribute and is useful before encoding categorical variables into numerical format for Machine Learning models.

## Output

The output is an array of unique values in the `education` column. For example: `['primary', 'secondary', 'tertiary', 'unknown']`, which shows all education levels present in the dataset.

```
[12]: array(['primary', 'secondary', 'tertiary', 'unknown'], dtype=object)
```

## 13 Encoding Education Attribute

```
[13]: df.education = df.education.map({
    'primary': 1,
    'secondary': 2,
    'tertiary': 3,
    'unknown': np.nan
})
df.head(10)
```

### Explanation

The `education` column contains categorical values representing the education level of customers. To prepare the data for Machine Learning models, these categorical values are converted to numerical values using the `map()` function. The mapping is as follows: '`primary`' = 1, '`secondary`' = 2, '`tertiary`' = 3, and '`unknown`' is replaced with `NaN` to indicate missing data. This transformation allows algorithms to process the education levels quantitatively.

### Output

The output displays the first ten rows of the dataset after the transformation. The `education` column now contains numerical values (1, 2, 3) corresponding to education levels, with `NaN` for unknown entries.

```
[13]:   age  job  marital  education  default  balance  housing  loan  contact \
0    30  0.0        1         1.0      no     1787       0     0  cellular
1    33  0.0        1         2.0      no     4789       1     1  cellular
2    35  1.0        0         3.0      no     1350       1     0  cellular
3    30  1.0        1         3.0      no     1476       1     1  unknown
4    59  0.0        1         2.0      no       0       1     0  unknown
5    35  1.0        0         3.0      no      747       0     0  cellular
6    36  0.0        1         3.0      no      307       1     0  cellular
7    39  1.0        1         2.0      no      147       1     0  cellular
8    41  1.0        1         3.0      no      221       1     0  unknown
9    43  0.0        1         1.0      no     -88       1     1  cellular

      day  month  duration  campaign  pdays  previous  poutcome  y
0     19      10       79         1      -1        0  unknown  no
1     11       5      220         1     339        4  failure  no
2     16       4      185         1     330        1  failure  no
3      3       6      199         4      -1        0  unknown  no
4      5       5      226         1      -1        0  unknown  no
5     23       2      141         2     176        3  failure  no
6     14       5      341         1     330        2  other   no
7      6       5      151         2      -1        0  unknown  no
8     14       5       57         2      -1        0  unknown  no
9     17       4      313         1     147        2  failure  no
```

## 14 Unique Values of the Outcome of Previous Marketing Campaign (poutcome)

```
[14]: df["poutcome"].unique()
```

### Explanation

The `unique()` function is used to identify all distinct values present in the `poutcome` column of the dataset. This column represents the result of the previous marketing campaign for each customer. Determining unique values helps in understanding the different categories present and assists in further preprocessing or encoding steps required for Machine Learning.

### Output

The output is an array of all unique values in the `poutcome` column, for example:

```
array(['unknown', 'failure', 'other', 'success'], dtype=object)
```

This shows that the column contains four distinct categories indicating the previous campaign outcome.

```
[14]: array(['unknown', 'failure', 'other', 'success'], dtype=object)
```

## 15 Encoding Poutcome Attribute

```
[15]: df.poutcome = df.poutcome.map({
    'unknown': np.nan,
    'failure': 1,
    'other': 2,
    'success': 3
})
df.head(10)
```

### Explanation

The `poutcome` column represents the outcome of the previous marketing campaign and is a categorical variable. To make it compatible with Machine Learning algorithms, it is mapped to numerical values using the `map()` function. The mapping assigns `failure` to 1, `other` to 2, `success` to 3, and replaces `unknown` values with `NaN` to handle missing information. This encoding simplifies the dataset while preserving meaningful distinctions between campaign outcomes.

### Output

The output displays the first ten rows of the dataset with the `poutcome` column transformed into numerical values. Entries that were `unknown` are now shown as `NaN`, while other outcomes are represented by 1, 2, or 3 accordingly.

```
[15]:   age  job  marital  education default  balance  housing  loan  contact \
0    30  0.0        1       1.0     no      1787        0      0  cellular
1    33  0.0        1       2.0     no      4789        1      1  cellular
2    35  1.0        0       3.0     no      1350        1      0  cellular
```

3	30	1.0	1	3.0	no	1476	1	1	unknown
4	59	0.0	1	2.0	no	0	1	0	unknown
5	35	1.0	0	3.0	no	747	0	0	cellular
6	36	0.0	1	3.0	no	307	1	0	cellular
7	39	1.0	1	2.0	no	147	1	0	cellular
8	41	1.0	1	3.0	no	221	1	0	unknown
9	43	0.0	1	1.0	no	-88	1	1	cellular
<hr/>									
0	19	NaN	79	1	-1	0	NaN	no	
1	11	NaN	220	1	339	4	1.0	no	
2	16	NaN	185	1	330	1	1.0	no	
3	3	NaN	199	4	-1	0	NaN	no	
4	5	NaN	226	1	-1	0	NaN	no	
5	23	NaN	141	2	176	3	1.0	no	
6	14	NaN	341	1	330	2	2.0	no	
7	6	NaN	151	2	-1	0	NaN	no	
8	14	NaN	57	2	-1	0	NaN	no	
9	17	NaN	313	1	147	2	1.0	no	

## 16 Normalizing the Balance Attribute

```
[16]: df["balance"] = df["balance"].apply(lambda v: (v - df["balance"].min())/
                                         (df["balance"].max() - df["balance"].min()))
df.head(10)
```

### Explanation

The balance column is a numerical feature representing the account balance of customers. To scale this feature between 0 and 1, min-max normalization is applied. The formula used is:

$$\text{normalized\_value} = \frac{v - \text{min(balance)}}{\text{max(balance)} - \text{min(balance)}}$$

where  $v$  is the original balance value. This transformation ensures that all values of balance lie within the range [0, 1], improving the performance and convergence of many Machine Learning algorithms.

### Output

The output displays the first ten rows of the dataset after normalization. The balance column values are now scaled between 0 and 1 while all other columns remain unchanged.

```
[16]:    age  job  marital  education default  balance  housing  loan  contact \
0   30  0.0      1       1.0     no  0.068455      0     0  cellular
1   33  0.0      1       2.0     no  0.108750      1     1  cellular
2   35  1.0      0       3.0     no  0.062590      1     0  cellular
3   30  1.0      1       3.0     no  0.064281      1     1  unknown
4   59  0.0      1       2.0     no  0.044469      1     0  unknown
5   35  1.0      0       3.0     no  0.054496      0     0  cellular
6   36  0.0      1       3.0     no  0.048590      1     0  cellular
7   39  1.0      1       2.0     no  0.046442      1     0  cellular
```

8	41	1.0	1	3.0	no	0.047436	1	0	unknown
9	43	0.0	1	1.0	no	0.043288	1	1	cellular
<hr/>									
0	19	NaN	79	1	-1	0	NaN	no	
1	11	NaN	220	1	339	4	1.0	no	
2	16	NaN	185	1	330	1	1.0	no	
3	3	NaN	199	4	-1	0	NaN	no	
4	5	NaN	226	1	-1	0	NaN	no	
5	23	NaN	141	2	176	3	1.0	no	
6	14	NaN	341	1	330	2	2.0	no	
7	6	NaN	151	2	-1	0	NaN	no	
8	14	NaN	57	2	-1	0	NaN	no	
9	17	NaN	313	1	147	2	1.0	no	

## 17 Normalization of pdays Attribute

```
[17]: df["pdays"] = df["pdays"].apply(lambda v: (v - df["pdays"].min())/
                                         (df["pdays"].max() - df["pdays"].min()))
df.head(10)
```

### Explanation

The pdays column, which represents the number of days since a client was last contacted, is normalized using the Min-Max scaling technique. This transformation scales all values to a range between 0 and 1, which helps improve the performance and convergence of Machine Learning algorithms.

### Output

The first ten rows of the dataset are displayed. The pdays column now contains values scaled between 0 and 1, while the other columns remain unchanged.

```
[17]:   age  job  marital  education default  balance  housing  loan  contact \
0    30  0.0      1        1.0     no  0.068455      0      0  cellular
1    33  0.0      1        2.0     no  0.108750      1      1  cellular
2    35  1.0      0        3.0     no  0.062590      1      0  cellular
3    30  1.0      1        3.0     no  0.064281      1      1  unknown
4    59  0.0      1        2.0     no  0.044469      1      0  unknown
5    35  1.0      0        3.0     no  0.054496      0      0  cellular
6    36  0.0      1        3.0     no  0.048590      1      0  cellular
7    39  1.0      1        2.0     no  0.046442      1      0  cellular
8    41  1.0      1        3.0     no  0.047436      1      0  unknown
9    43  0.0      1        1.0     no  0.043288      1      1  cellular

          day  month  duration  campaign  pdays  previous  poutcome  y
0     19    NaN       79         1  0.000000      0     NaN  no
1     11    NaN      220         1  0.389908      4     1.0  no
2     16    NaN      185         1  0.379587      1     1.0  no
3      3    NaN      199         4  0.000000      0     NaN  no
```

4	5	NaN	226	1	0.000000	0	NaN	no
5	23	NaN	141	2	0.202982	3	1.0	no
6	14	NaN	341	1	0.379587	2	2.0	no
7	6	NaN	151	2	0.000000	0	NaN	no
8	14	NaN	57	2	0.000000	0	NaN	no
9	17	NaN	313	1	0.169725	2	1.0	no

## 18 Feature Scaling using Min-Max Scaler

```
[18]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["duration"]] = scaler.fit_transform(df[["duration"]])
df[["pdays"]] = scaler.fit_transform(df[["pdays"]])
df.head()
```

### Explanation

Feature scaling is performed to normalize the numerical attributes duration and pdays into a fixed range, usually between 0 and 1, which helps in faster convergence and better performance of Machine Learning algorithms. The `MinMaxScaler` from `sklearn.preprocessing` is used. The `fit_transform()` method calculates the minimum and maximum values of each feature and scales all values accordingly. This ensures that the features contribute proportionally to the model training.

### Output

The output displays the first five rows of the dataset where the duration and pdays columns have been scaled to values between 0 and 1, while all other columns remain unchanged.

```
[18]:    age  job  marital  education  default  balance  housing  loan  contact \
0   30  0.0       1      1.0     no  0.068455      0      0  cellular
1   33  0.0       1      2.0     no  0.108750      1      1  cellular
2   35  1.0       0      3.0     no  0.062590      1      0  cellular
3   30  1.0       1      3.0     no  0.064281      1      1  unknown
4   59  0.0       1      2.0     no  0.044469      1      0  unknown

      day  month  duration  campaign  pdays  previous  poutcome  y
0    19    NaN  0.024826        1  0.000000      0    NaN  no
1    11    NaN  0.071500        1  0.389908      4    1.0  no
2    16    NaN  0.059914        1  0.379587      1    1.0  no
3     3    NaN  0.064548        4  0.000000      0    NaN  no
4     5    NaN  0.073486        1  0.000000      0    NaN  no
```